

PR_11.3 Dani Gayol Rodríguez

PR_11.3 Dani Gayol Rodríguez.....	1
Apartado A.....	1
Gestión de una base de datos en Amazon Redshift	2
Apartado B.....	2
Gestión de un esquema en una base de datos.....	3
Apartado C	3
Gestión de tablas en una base de datos.....	4
Apartado D	4
Gestión de vistas en una base de datos	5
Apartado E.....	5
Gestión de vistas materializadas en una base de datos.....	6
Apartado F.....	6
Gestión de procedimientos almacenados en una base de datos.....	6

Apartado A

Gestión de una base de datos en Amazon Redshift

Primero, vamos a conectarnos al cliente SQL de Amazon Redshift:

Edit connection for redshift-cluster-1

☐

IAM Identity Center

Connect to Amazon Redshift with your single sign-on credentials from your identity provider (IdP). Your cluster or workgroup must be enabled for IAM Identity Center.

☒

Other ways to connect [Learn more](#)

☐ Federated user

The principal tags of your IAM role or user must provide the connection details. You configure these tags in IAM or your identity provider (IdP).

☐ Temporary credentials using a database user name

The query editor v2 generates a temporary password to connect to the database.

☐ Temporary credentials using your IAM identity

The query editor v2 generates a temporary password to connect to the database as your IAM identity.

☒ Database user name and password

Provide a database user and password for the database that you are connecting to. The query editor v2 stores your credentials in AWS Secrets Manager on your behalf.

☐ AWS Secrets Manager

Choose a secret with credentials that are associated with the cluster or that you created in AWS Secrets Manager. Only secrets tagged with a key starting with 'Redshift' are listed.

Database

dev

The database name must be 1-64 characters. Valid characters are lowercase alphanumeric characters.

User name

awsuser


Password

.....


☐ Show password




Cancel

Save



Una vez conectados, ya podemos realizar este apartado:





Run

Limit 100

Explain

Isolated session

redshift-clust...

```
1 SELECT datname, datdba, datconnlimit
2 FROM pg_database_info
3 WHERE datname = 'qa';
```

Result 1 (1)

	datname	datdba	datconnlimit
	qa	100	50

Run

Limit 100

Explain

```
1 ALTER DATABASE qa RENAME TO prod;
```

Result 1

Summary

Returned rows: 0
Elapsed time: 271ms
Result set query:

```
/* RQEV2-BuvUiYP6GC */
ALTER DATABASE qa RENAME TO prod
```

Run

Limit 100

Explain

```
1 ALTER DATABASE prod OWNER TO awsuser;
```

Result 1

Summary

Returned rows: 0
Elapsed time: 301ms
Result set query:

```
/* RQEV2-nTNjXr5bMM */
ALTER DATABASE prod OWNER TO awsuser
```

```
▶ Run Limit 100 Explain Isolated session
```

```
1 ALTER DATABASE prod CONNECTION LIMIT 100;
```

Result 1

Summary

Returned rows: 0
Elapsed time: 243ms
Result set query:

```
/* RQEV2-KaRJWNbrvP */  
ALTER DATABASE prod CONNECTION LIMIT 100
```

```
▶ Run Limit 100 Explain Isolated session redshift-clust...
```

```
1 SELECT datname, datdba, datconlimit  
2 FROM pg_database_info  
3 WHERE datname = 'prod';
```

Result 1 (1)

	datname	datdba	datconlimit
	prod	100	100

Ahora, el último paso vamos a hacerlo desde “psql”, para ello hacemos lo siguiente:

```
Símbolo del sistema - \psql - x + v  
Microsoft Windows [Versión 10.0.26200.7840]  
(c) Microsoft Corporation. Todos los derechos reservados.  
C:\Users\Mañana>cd C:\Program Files\PostgreSQL\18\bin  
C:\Program Files\PostgreSQL\18\bin>set PGCLIENTENCODING=UTF8  
C:\Program Files\PostgreSQL\18\bin>.\psql -h "redshift-cluster-1.c7a12korjnpd.us-east-1.redshift.amazonaws.com" -d prod  
-p 5439 -U awsuser  
Contraseña para usuario awsuser:  
  
psql (18.3, servidor 8.0.2)  
ADVERTENCIA: psql versión mayor 18, servidor versión mayor 8.0.  
Algunas características de psql podrían no funcionar.  
ADVERTENCIA: El código de página de la consola (850) difiere del código  
de página de Windows (1252).  
Los caracteres de 8 bits pueden funcionar incorrectamente.  
Vea la página de referencia de psql «Notes for Windows users»  
para obtener más detalles.  
Conexión SSL (protocolo: TLSv1.3, cifrado: TLS_AES_128_GCM_SHA256, compresión: desactivado, ALPN: ninguno)  
Dígame «help» para obtener ayuda.  
prod=# |
```


```
prod=# SELECT procpid, username, datname FROM pg_stat_activity WHERE datname = 'prod';
ERROR:  syntax error at or near "clear" in context "clear", at line 1
LÍNEA 1: clear
        ^
prod=# |
```

```
prod=# SELECT pg_terminate_backend(procpid) FROM pg_stat_activity WHERE datname = 'prod' AND procpid <> pg_backend_pid();
pg_terminate_backend
-----
(0 filas)
```

```
dev=# DROP DATABASE prod;
DROP DATABASE
```

Apartado B

Gestión de un esquema en una base de datos



The screenshot shows a Redshift SQL editor interface. At the top, there is a toolbar with a 'Run' button, an upward arrow, and three toggle switches: 'Limit 100' (on), 'Explain' (off), and 'Isolated session' (on). A dropdown menu on the right shows 'redshift-c'. The SQL query is as follows:


```
1 create group finance_grp;
2 create group audit_grp;
3 create user finance_admin_usr with password 'Finance_admin_usr1';
```

Below the query, there are three tabs labeled 'Result 1', 'Result 2', and 'Result 3'. The 'Summary' section is visible, showing:

- Returned rows: 0
- Elapsed time: 297ms
- Result set query:

The result set query is displayed in a box:

```
/* RQEV2-0XMts4LU20 */
create user finance_admin_usr with password 'Finance_admin_usr1'
```



The screenshot shows a Redshift SQL editor interface. At the top, there is a toolbar with a 'Run' button, an upward arrow, and three toggle switches: 'Limit 100' (on), 'Explain' (off), and 'Isolated session' (on). A dropdown menu on the right shows 'redshift-c'. The SQL query is as follows:

```
1 CREATE schema finance authorization finance_admin_usr QUOTA 1 TB;
```

Below the query, there is a single tab labeled 'Result 1'. The 'Summary' section is visible, showing:

- Returned rows: 0
- Elapsed time: 253ms
- Result set query:

The result set query is displayed in a box:

```
/* RQEV2-XrKVQj0yLY */
CREATE schema finance authorization finance_admin_usr QUOTA 1 TB
```


▶ Run

⬆

Limit 100

Explain

Isolated session

1 GRANT USAGE on SCHEMA finance TO GROUP finance_grp;

☰ Result 1

Summary

Returned rows: 0

Elapsed time: 266ms

Result set query:

```
/* RQEV2-aDLPguXaok */
GRANT USAGE on SCHEMA finance TO GROUP finance_grp
```

▶ Run

⬆

Limit 100

Explain

Isolated session

1 GRANT USAGE on SCHEMA finance TO GROUP audit_grp;

☰ Result 1

Summary

Returned rows: 0

Elapsed time: 451ms

Result set query:

```
/* RQEV2-iMnGMfyMS0 */
GRANT USAGE on SCHEMA finance TO GROUP audit_grp
```

▶ Run ^

Limit 100

Explain

Isolated session

1 GRANT ALL ON schema finance to GROUP finance_grp;

Result 1

Summary

Returned rows: 0

Elapsed time: 351ms

Result set query:

/* RQEV2-5qs0P3Pm0m */

GRANT ALL ON schema finance to GROUP finance_grp

▶ Run ^

Limit 100

Explain

Isolated session

i

redshift

1 GRANT SELECT ON ALL TABLES IN SCHEMA finance TO GROUP audit_grp;

Result 1

Summary

Returned rows: 0

Elapsed time: 345ms

Result set query:

/* RQEV2-7mkPmphYdW */

GRANT SELECT ON ALL TABLES IN SCHEMA finance TO GROUP audit_grp

Ahora para verificar hacemos lo siguiente:

Run Limit 100 Explain Isolated session

```
1 SELECT nsname as schema, username as owner
2 FROM pg_namespace, pg_user
3 WHERE pg_namespace.nspowner = pg_user.usename
4 and pg_namespace.nspname = 'finance';
```

Result 1 (1)

schema	owner
finance	finance_admin_usr

Run Limit 100 Explain Isolated session

```
1 CREATE TABLE finance.foo (bar int);
```

Result 1

Summary

Returned rows: 0
Elapsed time: 360ms
Result set query:

```
/* RQEV2-JSm4a5oWZR */
CREATE TABLE finance.foo (bar int)
```

Run Limit 100 Explain Isolated session

```
1 SELECT * FROM finance.foo;
```

Result 1

bar

▶ Run

⬆

Limit 100

Explain

Isolated

1 set search_path to '\$user', finance, public;

Result 1

Summary

Returned rows: 0

Query ID: 220793

Elapsed time: 8ms

Result set query:

```
/* RQEV2-kCuRF3dNxp */
set search_path to '$user', finance, public
```

▶ Run

⬆

Limit 100


1 SELECT * FROM foo;

Result 1

bar

Apartado C

Gestión de tablas en una base de datos



The screenshot shows a SQL IDE interface. At the top, there is a toolbar with a 'Run' button (a play icon), a 'Limit 100' toggle switch (which is turned on), and an 'Explain' button. Below the toolbar, the SQL editor contains a query to create a table named 'customer' in the 'finance' schema. The query is as follows:

```
1 CREATE TABLE finance.customer (  
2     customer_number INTEGER,  
3     first_name VARCHAR(50),  
4     last_name VARCHAR(50),  
5     date_of_birth DATE  
6 );
```

Below the editor, there is a tab labeled 'Result 1'. The results pane shows a 'Summary' section with the following information:

- Returned rows: 0
- Query ID: 220807
- Elapsed time: 341ms
- Result set query:

The 'Result set query' section displays the executed SQL statement, which is identical to the one in the editor, preceded by a comment: `/* RQEV2-mAHfs3WFua */`.

▶ Run ^

Limit 100

Explain

Isolated session

1

2

3

4

5

6

7

8

9

10

11

insert into finance. customer values

(1, 'foo', 'bar', '1980-01-01'),

(2, 'john', 'smith', '1990-12-01'),

(3, 'spock', 'spock', '1970-12-01'),

(4, 'scotty', 'scotty', '1975-02-01'),

(5, 'seven', 'of nine', '1990-04-01'),

(6, 'kathryn', 'janeway', '1995-07-01'),

(7, 'tuvok', 'tuvok', '1960-06-10'),

(8, 'john', 'smith', '1965-12-01'),

(9, 'The Doctor', 'The Doctor', '1979-12-01'),

(10, 'B Elana', 'Torres', '2000-08-01');

Result 1

Summary

Affected rows: 10

Query ID: 220890

Elapsed time: 684ms

Result set query:

/* RQEV2-EGIIjMEHrf */

insert into finance. customer values

(1, 'foo', 'bar', '1980-01-01'),

(2, 'john', 'smith', '1990-12-01'),

(3, 'spock', 'spock', '1970-12-01'),

(4, 'scotty', 'scotty', '1975-02-01'),

(5, 'seven', 'of nine', '1990-04-01'),

(6, 'kathryn', 'janeway', '1995-07-01'),

(7, 'tuvok', 'tuvok', '1960-06-10'),

(8, 'john', 'smith', '1965-12-01'),

(9, 'The Doctor', 'The Doctor', '1979-12-01'),

(10, 'B Elana', 'Torres', '2000-08-01')

▶ Run ^

Limit 100

Explain

Isolated session

redshift-clust...

dev

1

2

3

4

SELECT "schema", table_id, "table", encoded, diststyle, sortkey1, size, tbl_rows

FROM svv_Table_info

WHERE "table" = 'customer'

and "schema" = 'finance';

Result 1 (1)

☐	schema	table_id	table	encoded	diststyle	sortkey1
☐	finance	112717	customer	Y, AUTO(ENCODE)	AUTO(ALL)	AUTO(SORTKEY)

Run

Limit 100

Explain

Isolated session

redshift-clust...

dev

```
1 SELECT *
2 FROM finance.customer
3 WHERE EXTRACT(year from date_of_birth) < 1980;
```

Result 1 (5)

	customer number	first name	last name	date of birth
☐	3	spock	spock	1970-12-01
☐	4	scotty	scotty	1975-02-01
☐	7	tuvok	tuvok	1960-06-10
☐	8	john	smith	1965-12-01
☐	9	The Doctor	The Doctor	1979-12-01

Run

Limit 100

Explain

Isolated sess

```
1 CREATE TABLE finance.customer_dob_1980 as
2 SELECT *
3 FROM finance.customer
4 WHERE EXTRACT(year from date_of_birth) = 1980 ;
```

Result 1

Summary

Returned rows: 0

Query ID: 220960

Elapsed time: 991ms

Result set query:

```
/* RQEV2-UIucdRip2N */
CREATE TABLE finance.customer_dob_1980 as
SELECT *
FROM finance.customer
WHERE EXTRACT(year from date_of_birth) = 1980
```

Run

Limit 100

Explain

Isolated session

redshift-clust...

dev

```
1 CREATE temporary TABLE #customer(custid integer IDENTITY(1,1), customer_number integer);
2 INSERT INTO #customer (customer_number) values(1);
3 SELECT * FROM #customer;
```

Result 1

Result 2

Result 3 (1)

	custid	customer number
☐	1	1

Ahora si intentamos hacerlo desde el cliente SQL, nos da error de que no existe ya que es una tabla temporal y solo se muestra en la sesión actual:

```
dev=# SELECT * FROM #customer;  
ERROR:  relation "#customer" does not exist  
dev=#
```


Apartado D

Gestión de vistas en una base de datos

The screenshot shows a SQL query execution interface. At the top, there are controls for running the query, including a 'Run' button, a 'Limit 100' toggle, an 'Explain' toggle, an 'Isolated session' toggle, and dropdown menus for the cluster ('redshift-clust...') and the database ('dev'). The query is as follows:

```
1 CREATE VIEW finance.customer_vw
2 AS
3 SELECT customer_number, first_name, last_name, EXTRACT(year FROM date_of_birth) AS year_of_birth
4 FROM finance.customer;
```

Below the query, the 'Result 1' tab is selected, showing a 'Summary' section. The summary indicates that 0 rows were returned, the query ID is 409291, and the elapsed time is 295ms. The 'Result set query' section shows the same query that was executed, including a comment: `/* RQEV2-EwUq8GN4Yy */`.

The screenshot shows a SQL query execution interface. At the top, there are controls for running the query, including a 'Run' button, a 'Limit 100' toggle, an 'Explain' toggle, an 'Isolated session' toggle, and dropdown menus for the cluster ('redshift-clust...') and the database ('dev'). The query is as follows:

```
1 CREATE VIEW finance.customer_lbv
2 AS
3 SELECT customer_number, first_name, last_name, EXTRACT(year FROM date_of_birth) AS year_of_birth
4 FROM finance.customer
5 WITH no schema binding;
```

Below the query, the 'Result 1' tab is selected, showing a 'Summary' section. The summary indicates that 0 rows were returned, the query ID is 409291, and the elapsed time is 242ms. The 'Result set query' section shows the same query that was executed, including a comment: `/* RQEV2-ZmbfGjDCip */`.

Ahora vamos a verificar que se creó correctamente:

Run Limit 100 Explain Isolated session redshift-clust...

```

1 SELECT table_schema as schema_name, table_name as view_name, view_definition
2 FROM information_schema.views
3 WHERE table_schema not in ('information_schema', 'pg_catalog' )
4 ORDER by schema_name, view_name;

```

Result 1 (3)

	schema_name	view_name	view_definition
☐	finance	customer_lbvw	/* RQEV2-ZmbfGjDCip */
☐	finance	customer_vw	SELECT customer.custo...
☐	pg_auto_copy	copy_job_detail	NULL

Run Limit 100 Explain Isolated session redshift-clust... dev

```

1 SELECT *
2 FROM finance.customer_vw LIMIT 5;

```

Result 1 (5)

	customer number	first name	last name	year of birth
☐	1	foo	bar	1980
☐	2	john	smith	1990
☐	3	spock	spock	1970
☐	4	scotty	scotty	1975
☐	5	seven	of nine	1990

Run Limit 100 Explain

```

1 DROP VIEW finance.customer_vw;

```

Result 1

Summary

Returned rows: 0
 Query ID: 409366
 Elapsed time: 312ms
 Result set query:

```

/* RQEV2-NvWRg4a9wP */
DROP VIEW finance.customer_vw

```


Apartado E

Gestión de vistas materializadas en una base de datos

▶ Run ^

Limit 100

Explain

Isolated session

redshift-clust...

```
1 CREATE MATERIALIZED VIEW finance.customer_agg_mv
2 AUTO REFRESH YES
3 AS
4 SELECT EXTRACT(year FROM date_of_birth) AS year_of_birth, count(1) customer_cnt
5 FROM finance.customer
6 group by EXTRACT(year FROM date_of_birth);
```

Result 1

Summary

Returned rows: 0
Query ID: 409401
Elapsed time: 911ms
Result set query:

```
/* RQEV2-vvgSicPLrU */
CREATE MATERIALIZED VIEW finance.customer_agg_mv
AUTO REFRESH YES
AS
SELECT EXTRACT(year FROM date_of_birth) AS year_of_birth, count(1) customer_cnt
FROM finance.customer
group by EXTRACT(year FROM date_of_birth)
```

▶ Run ^

Limit 100

Explain

Isolated session

```
1 SELECT *
2 FROM finance.customer_agg_mv
3 LIMIT 5;
```

Result 1 (5)

☐	year of birth	customer cnt
☐	1980	1
☐	1970	1
☐	1975	1
☐	1960	1
☐	1965	1

Run Limit 100 Explain Isolated session redshift-clust... dev Schedule								
<pre> 1 SELECT * 2 FROM SVV_MV_INFO 3 WHERE name = 'customer_agg_mv'; </pre>								
Result 1 (1) Export Chart								
database name	schema name	user name	name	is stale	state	autorewrite	autorefresh	
dev	finance	awsuser	customer_agg_mv	f	1	t	t	

Run Limit 100 Explain Isolated session								
<pre> 1 INSERT INTO finance.customer VALUES 2 (11, 'mark', 'bar', '1980-02-01'), 3 (12, 'pete', 'smith', '1990-2-01'), 4 (13, 'woofy', 'spock', '1980-11-01'), 5 (14, 'woofy jr', 'scotty', '1975-03-01'), 6 (15, 'eleven', 'of nine', '1990-07-01'); </pre>								
Result 1								
Summary Affected rows: 5 Query ID: 409440 Elapsed time: 655ms Result set query: <pre> /* RQEV2-h6VZZP0cn3 */ INSERT INTO finance.customer VALUES (11, 'mark', 'bar', '1980-02-01'), (12, 'pete', 'smith', '1990-2-01'), (13, 'woofy', 'spock', '1980-11-01'), (14, 'woofy jr', 'scotty', '1975-03-01'), (15, 'eleven', 'of nine', '1990-07-01') </pre>								

Run Limit 100 Explain Isolated session redshift-clust...								
<pre> 1 SELECT name, is_stale, state 2 FROM SVV_MV_INFO 3 WHERE name = 'customer_agg_mv'; </pre>								
Result 1 (1)								
name	is stale	state						
customer_agg_mv	f	1						

▶ Run

Limit 100

Explain

Isolated session

1 REFRESH MATERIALIZED VIEW finance.customer_agg_mv;

Result 1

Summary

Info:

- Materialized view customer_agg_mv is already up to date.

Returned rows: 0

Query ID: 409467

Elapsed time: 260ms

Result set query:

```
/* RQEV2-ZZDVxLfQgI */
REFRESH MATERIALIZED VIEW finance.customer_agg_mv
```

▶ Run

Limit 100

Explain

Isolated session

redshift-clust...

1 SELECT name, is_stale, state

2 FROM SVV_MV_INFO


3 WHERE name = 'customer_agg_mv';

Result 1 (1)

☐	name	is stale	state
☐	customer_agg_mv	f	1

Apartado F

Gestión de procedimientos almacenados en una base de datos



▶ Run ^ Limit 100 Explain Isolated session redshift-clust... dev

```
1 CREATE OR REPLACE procedure sp_cookbook(indate in date, records_out INOUT refcursor) as $$ declare integer_var int;
2 BEGIN
3     RAISE INFO 'running first cookbook storedprocedure on date %', indate;
4     DROP TABLE if exists cookbook.cookbook_tbl;
5     CREATE TABLE cookbook.cookbook_tbl(recipe_name varchar(50), recipe_date date);
6
7     INSERT INTO cookbook.cookbook_tbl values('stored procedure', indate);
8     GET DIAGNOSTICS integer_var := ROW_COUNT;
9     RAISE INFO 'rows inserted into cookbook_tbl = %', integer_var;
10    OPEN records_out FOR SELECT * FROM cookbook. cookbook_tbl;
11 END;
12 $$ LANGUAGE plpgsql;
```

Result 1

Summary

Returned rows: 0
Query ID: 409498
Elapsed time: 253ms
Result set query:

```
/* RQEV2-BQtjwMEHrn */
CREATE OR REPLACE procedure sp_cookbook(indate in date, records_out INOUT refcursor) as $$ declare integer_var int;
BEGIN
    RAISE INFO 'running first cookbook storedprocedure on date %', indate;
    DROP TABLE if exists cookbook.cookbook_tbl;
    CREATE TABLE cookbook.cookbook_tbl(recipe_name varchar(50), recipe_date date);

    INSERT INTO cookbook.cookbook_tbl values('stored procedure', indate);
    GET DIAGNOSTICS integer_var := ROW_COUNT;
    RAISE INFO 'rows inserted into cookbook_tbl = %', integer_var;
    OPEN records_out FOR SELECT * FROM cookbook. cookbook_tbl;
END;
$$ LANGUAGE plpgsql
```

Al hacer la siguiente operación, nos dará un error ya que nosotros tenemos un Cluster “mono nodo”:

▶ Run

⌵

Limit 100

⌵

Explain

⌵

Isolated ses

⌵

1 CALL sp_cookbook(current_date, 'inputcursor');

2 FETCH ALL FROM inputcursor;

Result 1

Result 2

Summary

❗ ERROR: Fetch ALL is not supported on single-node clusters.

Query ID: 409891

Elapsed time: 0 seconds

Result set query:

/* RQEV2-9AafcptJNo */

FETCH ALL FROM inputcursor

▶ Run

⌵

Limit 100

⌵

Explain

⌵

Isolated session

⌵

redshift-clust...

ⓘ

1 SHOW PROCEDURE sp_cookbook(indate in date, records_out INOUT refcursor);

Result 1 (1)

☐

Stored Procedure Def...

☐

CREATE OR REPLACE ...

▶ Run

Limit 100

Explain

Isolated session

redshift-clust...

dev

```
1 CREATE OR REPLACE procedure public.sp_self_service(tblName in varchar(60)) as $$
2 BEGIN
3     RAISE INFO 'running sp_self_service to drop table %', tblName;
4     execute 'drop table if exists cookbook.' || tblName;
5     RAISE INFO 'table dropped %', tblName;
6 END;
7 $$ LANGUAGE plpgsql
8 SECURITY DEFINER;
```

Result 1

Summary

Returned rows: 0

Query ID: 410145

Elapsed time: 370ms

Result set query:

```
/* RQEV2-1Dzc92Cn0j */
CREATE OR REPLACE procedure public.sp_self_service(tblName in varchar(60)) as $$
BEGIN
    RAISE INFO 'running sp_self_service to drop table %', tblName;
    execute 'drop table if exists cookbook.' || tblName;
    RAISE INFO 'table dropped %', tblName;
END;
$$ LANGUAGE plpgsql
SECURITY DEFINER
```

▶ Run



Limit 100



Explain



Isolated session



redshift-clust...

dev

```
1 CREATE USER user1 WITH password 'Cookbook1';
2
3 -- Otorga a user1 permiso para ejecutar el procedimiento sp_self_service.
4 GRANT execute on procedure public.sp_self_service(tblName in varchar(60)) to user1;
5
6 -- Cambia el usuario activo de la sesión a user1.
7 -- A partir de aquí, todo se ejecuta como si fueras user1.
8 SET SESSION authorization user1;
9 SELECT current_user;
10
11 -- Si user1 no tiene permisos sobre esa tabla,
12 -- Esta sentencia fallará con error de permisos.
13 DROP TABLE cookbook.cookbook_tbl;
```



Result 1



Result 2



Result 3



Result 4 (1)



Result 5

Summary

❗ ERROR: permission denied for schema cookbook [ErrorId: 1-69b7c5e7-59f7dceb1cd25b097e341035]

Query ID: 409891

Elapsed time: 0 seconds

Result set query:

```
/* RQEV2-ojDnVfWGuv */
-- Si user1 no tiene permisos sobre esa tabla,
-- Esta sentencia fallará con error de permisos.
DROP TABLE cookbook.cookbook_tbl
```

▶ Run

⌵

Limit 100

Explain

Isolated

1

SET SESSION authorization user1;

2

SELECT current_user;

3

CALL public.sp_self_service('cookbook_tbl');

Result 1

Result 2 (1)

Result 3

Summary

Info:

- running sp_self_service to drop table cookbook_tbl
- table dropped cookbook_tbl

Returned rows: 0

Query ID: 410225

Elapsed time: 283ms

Result set query:

/* RQEV2-Kk4gMKtJPO */

CALL public.sp_self_service('cookbook_tbl')